



Artificial Intelligence in K-12 Public Schools

Maryland's Framework for Responsible AI in Education

Office of Teaching and Learning | Division of College and Career Pathways

Last updated: June 2026

Table of Contents

Purpose and Executive Summary	2
Vision for Artificial Intelligence in Maryland Schools	2
Scope and Definitions	3
Roles and Responsibilities.....	4
Data Privacy and Security	5
Equity and Access	6
Academic Integrity and Assessment	7
Professional Learning.....	8
Student-Centered Learning with AI	9
Student AI Literacy and Curriculum Integration.....	10
Operational and Administrative Uses.....	11
Procurement and Vendor Governance	12
Conclusion	13
References.....	14

Purpose and Executive Summary

Artificial intelligence is increasingly shaping the tools used in K-12 education. It can personalize learning, support educators, improve operations, and help students develop skills needed for a changing economy. At the same time, AI introduces risks related to privacy, bias, equity, academic integrity, and over-reliance on automated systems.

This guidance establishes Maryland's statewide framework for the responsible use of artificial intelligence in K-12 education. It supports local education agencies, schools, and partners in making informed decisions about if and how AI is used in teaching, learning, and school operations, while affirming that technology must serve educational goals and that human judgment, relationships, and student wellbeing remain central.

Maryland's approach aligns with the Blueprint for Maryland's Future, statewide computer science and digital literacy efforts, and broader state and national work on safe, ethical AI. AI is treated not as a separate initiative, but as a set of tools embedded across curriculum, instruction, assessment, data governance, and workforce development.

The guidance sets a clear vision and principles for AI use, establishes expectations and guardrails for implementation, and provides a framework for governance and continuous improvement as technology and evidence evolve. It is intended for state leaders, LEA and school administrators, educators, students and families, and vendors, balancing statewide expectations with local flexibility.

Vision for Artificial Intelligence in Maryland Schools

Maryland's vision for artificial intelligence in education is grounded in a commitment to rigorous, equitable learning for every student, supported by educators who use emerging technologies responsibly. When implemented with care, AI can deepen learning, support instruction, and strengthen school operations, while human relationships, professional expertise, and student wellbeing remain central.

Artificial intelligence should enhance, not replace, the educators, counselors, and school leaders who remain Maryland's primary decision makers. While AI may assist with instruction, feedback, or administrative tasks, it cannot replace professional judgment, cultural competence, or the trust built through human relationships.

Maryland's vision also emphasizes preparing students to understand and critically evaluate AI systems that influence daily life. AI literacy is a core skill, enabling learners to assess reliability, identify bias, understand underlying assumptions, and make informed choices about technology use. Equity, transparency, privacy, and accountability are foundational. Maryland is committed to ensuring equitable access to AI-supported learning opportunities and to protecting student data through clear communication, human oversight, and compliance with all state and federal privacy, cybersecurity, and ethical standards.

Recognizing that AI will continue to evolve, Maryland's approach balances stability with adaptability. This guidance will be reviewed and updated regularly to reflect research, implementation experience,

and feedback, ensuring AI strengthens educational opportunities, protects human dignity, and prepares all students for success in a complex technology-driven future.

Scope and Definitions

This guidance applies to the adoption, use, evaluation, and governance of artificial intelligence in Maryland public schools. It establishes statewide expectations for responsible AI use by students, educators, school and district leaders, and education partners, and defines key terms, scope, and boundaries to support consistent local decision making.

DEFINITIONS TO KNOW

Artificial intelligence (AI) refers to computational systems that perform tasks associated with human intelligence, including pattern recognition, content generation, prediction, and adaptive output.

Machine learning is a subset of AI in which systems improve performance by identifying patterns in data.

Generative AI produces new content such as text, images, audio, video, or code in response to user input.

Large language models (LLM) are a type of generative AI that produce text based on statistical patterns in training data and do not possess understanding or intent.

AI literacy refers to the knowledge and skills needed to understand, evaluate, and use AI responsibly, including awareness of limitations, bias, and appropriate use.

Assistive AI tools support accessibility and instructional access, instructional AI tools support teaching and learning, and operational AI tools support school and district functions.

SCOPE

The scope of this guidance includes AI use for teaching and learning, student AI literacy and curricular integration, professional learning and educator practice, data privacy and ethical safeguards, procurement and evaluation, operational uses, and governance, monitoring, and accountability. It applies to all AI tools that interact with Maryland students, staff, or data, regardless of how they are accessed or provided.

This document establishes statewide expectations but does not approve or prohibit specific tools, require adoption of particular technologies, replace existing law, serve as a technical specification, or apply outside K-12 public education. Local education agencies retain authority to select tools and design implementation strategies aligned with this guidance.

Roles and Responsibilities

Responsible use of artificial intelligence in Maryland schools requires coordinated governance across the education system. AI oversight cannot rest with individual educators or isolated departments; it depends on shared responsibility, clear roles, and transparent processes among state agencies, districts, schools, educators, students, families, and vendors (Holmes et al., 2021; Floridi & Cows, 2019).

	<u>MSDE</u> provides statewide leadership by establishing AI policy expectations and guardrails related to equity, privacy, security, accessibility, instructional quality, and civil rights compliance. MSDE offers guidance, resources, and professional learning, monitors statewide trends, and updates policies as technology and evidence evolve. MSDE does not approve or endorse specific AI tools; tool selection and oversight remain local responsibilities.
	<u>Local education agencies</u> implement statewide expectations by adopting district AI policies, vetting, and approving tools, ensuring legal and privacy compliance, providing professional learning, monitoring systems for bias or unintended impacts, maintaining cybersecurity, communicating with families, and reporting concerns to MSDE.
	<u>School administrators</u> oversee day-to-day implementation, establish expectations for staff and students, support educators, address misuse or harm, communicate with families, and ensure AI use does not produce discriminatory outcomes.
	<u>Educators</u> are the primary decision makers for instructional AI use. They are responsible for ethical and critical use, protecting academic integrity and student privacy, modeling responsible practices, and reporting concerns.
	<u>Students</u> are responsible for following AI-use expectations, maintaining academic integrity, acknowledging permitted AI assistance, evaluating outputs critically, protecting their privacy, and reporting misuse.
	<u>Families</u> are partners in responsible AI use and are expected to review school communications, support appropriate practices at home, raise questions, and report concerns.
	<u>Vendors</u> must meet Maryland's standards for transparency, privacy, accessibility, equity, and security, including accurate disclosures, legal compliance, fairness testing, notification of system changes, and participation in audits. Tools that fail to meet these standards must not be used.

AI governance in Maryland is an ongoing, shared process grounded in accountability, transparency, and continuous improvement to ensure AI supports equitable opportunity and high-quality learning statewide.

Data Privacy and Security

Artificial intelligence increases the complexity of data privacy and security in Maryland schools because these systems rely on large volumes of data and often operate in opaque ways. Research shows that AI amplifies existing risks related to data collection, surveillance, re-identification, and cybersecurity, particularly for minors (Crawford, 2021; Reddy et al., 2020). Maryland's expectation is clear: AI use must never compromise student privacy or weaken data protections.

AI-related privacy risks include the collection of sensitive information through user inputs, the possibility of re-identifying anonymized data, expanded cybersecurity vulnerabilities, and limited transparency about how vendors store, use, or share data. These risks require proactive governance, strong safeguards, and continuous oversight.

PRIVACY & CIVIL RIGHTS LAWS

All AI use in Maryland schools must comply with federal and state privacy and civil rights laws, including FERPA, COPPA, PPRa, and applicable Maryland regulations. Algorithm-assisted decisions remain subject to nondiscrimination requirements, and districts retain full responsibility for protecting student rights (U.S. Department of Education, 2024). Schools must be transparent with families about how AI tools handle student data.

DISTRICT EXPECTATIONS

Maryland expects districts to minimize data collection by limiting student data shared with AI systems, avoiding entry of personally identifiable information into generative tools unless explicitly approved, restricting unnecessary data access, and providing alternatives when AI use is inappropriate. Collection of biometric, behavioral, or location data is prohibited unless legally required and educationally justified.

AI adoption increases cybersecurity risk. Districts must implement robust system protections, including secure authentication, encryption, vendor security requirements, staff training, and tested incident response plans. Privacy protections cannot be upheld without robust cybersecurity.

AI systems must operate with documented human oversight. Any AI-supported decision affecting students requires human review, clear procedures for addressing errors or harm, and regular audits for bias, accuracy, and unintended data collection.

Transparency with families is essential. Districts must clearly communicate what AI tools are used, what data are collected, how long data are retained, and when opt-out options apply, using accessible and culturally responsive formats.

Maryland will support districts through statewide guidance, templates, training, and reporting structures. The State's priority is to ensure that all AI used in public schools protects student privacy, strengthens security, and maintains public trust.

VENDOR EXPECTATIONS

AI vendors must meet strict privacy, security, and transparency standards. Districts must ensure contractual limits on data use, retention, and deletion; prohibit advertising, profiling, or unrelated model training; require disclosure of data practices; and ensure compliance with breach notification laws. Vendors that cannot demonstrate secure, education-focused data practices must not be used.

Equity and Access

Maryland is committed to ensuring that artificial intelligence strengthens, rather than undermines, equitable access to high-quality learning and support services. While AI can expand opportunity, research shows that algorithmic systems often reflect and reinforce existing racial, linguistic, economic, and disability-based inequities if not carefully designed and monitored (Benjamin, 2019; Noble, 2018). This guidance establishes expectations to ensure that AI use in Maryland schools promotes fairness, protects civil rights, and supports diverse learners.

BE AWARE OF BIAS

AI systems are shaped by their data and design choices and can produce biased or inequitable outcomes without oversight. Risks include biased outputs, unequal access to devices and connectivity, uneven instructional access to AI-supported learning, over-reliance on automation for marginalized students, and opaque decision systems that reproduce inequities through scheduling, early warning, or resource allocation tools. Automation bias further increases the risk of uncritical reliance on AI-generated recommendations (Goddard et al., 2012).

LEA EXPECTATIONS

Maryland expects LEAs to implement AI in ways that expand equitable access to opportunity. This includes ensuring reliable access to devices and connectivity, selecting high-quality tools that meet accessibility and language needs, designing instruction that maintains rigorous expectations, monitoring for biased or harmful outputs, prioritizing support for historically underserved communities, and communicating transparently with families in culturally and linguistically responsive ways (CAST, 2018; Koehn & Knowles, 2017; Reich et al., 2020).

All AI use must comply with federal and state civil rights laws. Algorithm-assisted decisions remain subject to nondiscrimination requirements, and schools retain responsibility for preventing discriminatory outcomes (U.S. Department of Education, 2024). AI may support accessibility but cannot replace legally required evaluations, services, accommodations, or language access under IDEA, Section 504, the ADA, or related statutes. Any AI system affecting student opportunities must include documented human oversight.

Maryland will monitor equity impacts of AI use statewide and expects LEAs to do the same within existing equity review structures. Where disparities in access, outcomes, or treatment emerge, corrective action is required.

The goal is not only to prevent harm, but to use AI deliberately to expand opportunity, support rigorous learning, strengthen family engagement, and connect students to future pathways. Achieving this requires sustained attention, transparency, and a clear commitment to ensuring that AI contributes positively to the educational experiences of all Maryland students.

Academic Integrity and Assessment

Artificial intelligence challenges traditional approaches to academic integrity and assessment by enabling students to generate work that may obscure original thinking. Maryland schools must establish clear expectations that protect the authenticity of student work, support responsible AI use, and ensure assessments measure student understanding.

Academic dishonesty occurs when students misrepresent AI-generated content as their own or use AI to bypass intended learning. Schools must clearly define when AI use is permitted or prohibited, when unaided work is required, and how AI assistance must be acknowledged.

Maryland does not endorse AI-detection tools as a primary means of enforcing academic integrity, as research shows they are unreliable and inequitable (Liang et al., 2023). Integrity should instead be supported through instructional design, transparency, and human judgment.

Assessment practices must evolve to emphasize reasoning, analysis, reflection, and process over recall. Tasks should require visible evidence of student thinking and reduce opportunities for AI substitution.

When AI use is allowed, students must document its role. Any use that obscures authorship, fabricates information, or replaces essential skill development violates academic integrity.

Students must learn to critically evaluate AI outputs, verify information, and understand AI limitations. Educators require support to design AI-resilient assessments and address integrity concerns through instruction rather than punishment.

Maryland's goal is to ensure that academic work reflects authentic learning, ethical participation, and meaningful intellectual growth in an AI-enabled environment.



Professional Learning

Artificial intelligence can strengthen teaching and learning only when educators are prepared, supported, and empowered to use it responsibly. Research shows that educator understanding, confidence, and professional judgment—not technology—determine whether AI improves outcomes (Zhai et al., 2023). Maryland's guidance, therefore, centers educator preparation and ongoing professional learning as essential to effective AI use.

HOW AI CAN SUPPORT INSTRUCTION

AI is not a substitute for teacher expertise. It may support instruction by reducing routine workload, expanding access to materials, or offering feedback, but educators remain the final decision makers. Without explicit preparation, educators may over-rely on AI or avoid it entirely, increasing risks related to bias, misinformation, and misuse (Holmes et al., 2021). Professional learning must therefore emphasize critical evaluation, verification, and reflective practice, particularly given the documented tendency to over-trust automated recommendations (Goddard et al., 2012).

AI-focused professional learning equips educators to use these tools safely, ethically, and effectively in student-centered ways. Core expectations include understanding foundational AI concepts, evaluating outputs for accuracy and bias, designing assignments that preserve academic integrity, communicating transparently with students and families, protecting student data, and recognizing when AI use is inappropriate. These priorities align with national guidance emphasizing that AI should strengthen, not diminish, human expertise (U.S. Department of Education, 2023).

PROFESSIONAL DEVELOPMENT

Professional learning must be role specific. Teachers require classroom-focused preparation tied to discipline-specific pedagogy. Administrators need understanding of policy, implementation, and risk management. District leaders require expertise in procurement, vendor oversight, equity monitoring, and governance. Specialists and support personnel need targeted guidance related to civil rights protections, service determination, and risks of bias in automated systems, consistent with federal requirements (U.S. Department of Education, OCR & OET, 2024).

Maryland will adopt core educator AI competencies aligned with emerging national and international frameworks (UNESCO, 2023). Educators are expected to understand what AI can and cannot do, evaluate outputs critically, integrate AI in ways that strengthen student thinking, teach responsible student use, protect privacy, monitor impacts on equity and learning, and report concerns related to bias or misuse.

Effective professional learning must be ongoing, job-embedded, and collaborative. Maryland will support sustained learning, coaching, and modeling, competency-based pathways, communities of practice, and just-in-time resources through the MSDE AI Hub (Holmes et al., 2019). Professional learning will also emphasize healthy skepticism, including identifying biased or fabricated outputs and understanding the limitations of generative models.

AI may help reduce administrative burden and support educator well-being, but all AI-generated materials require human review for accuracy, cultural responsiveness, and instructional alignment. Maryland will continue to update professional learning, competencies, and supports as research and tools evolve, ensuring that all educators have the knowledge and resources needed to use AI safely, ethically, and effectively in service of student learning.



Student-Centered Learning with AI

Artificial intelligence can enhance learning when it supports student agency, deepens understanding, and reinforces the vital role of educators. Used responsibly, AI may help differentiate instruction, support multilingual learners, provide feedback, and reduce administrative burden, while preserving academic integrity and teacher expertise. This guidance establishes expectations for student-centered AI use that complements, rather than replaces, student thinking.

COMPLIMENTING STUDENT EDUCATION

AI should be used to support engagement with complex ideas, provide appropriate scaffolding and feedback, and personalize learning in ways that respect student identity and experience. It must not be used to bypass learning goals or substitute for essential skill development. Educators determine when and how AI supports instruction, and student use must always occur within structures that promote original thinking and academic honesty.

Instructional Activity	Learning tasks should be intentionally designed so that student reasoning, analysis, reflection, and decision making are essential and visible. Assignments that can be completed primarily by AI offer limited instructional value and undermine skill development. Effective tasks require evidence-based reasoning, discipline-specific thinking, personal reflection, and iterative work that reveals learning over time.
Digital Citizenship	Students should learn to use AI transparently and responsibly. When AI use is permitted, students must acknowledge its role and be able to explain how outputs were evaluated, modified, or integrated into their own work. Misrepresenting AI-generated content as original work constitutes a violation of academic integrity. Documentation of AI use is intended to support learning, not punishment.
Accessibility	AI may support accessibility and multilingual learning through translation, varied reading levels, captioning, and assistive technologies. These tools must be used in compliance with federal and state civil rights laws and cannot replace required accommodations, specialized instruction, or professional judgment. Decisions about AI use for students with disabilities or multilingual learners should be made collaboratively with appropriate educators and families.

EDUCATOR EXPECTATIONS	LEA EXPECTATIONS
Educators guide all student use of AI. They establish expectations, model critical evaluation, monitor impacts on equity and learning, and ensure that AI does not overshadow disciplinary knowledge or relationships. AI functions as a tool within the learning environment, not a substitute for instructional expertise.	Student use of AI must also support safe and ethical learning environments. Schools should set clear expectations related to data privacy, bias awareness, digital citizenship, and reporting concerns. These practices help students develop responsible habits for engaging with intelligent technologies beyond school.



Student AI Literacy and Curriculum Integration

Artificial intelligence increasingly shapes how students learn, communicate, and participate in life. Preparing Maryland students for this reality requires intentional, developmentally appropriate instruction in AI literacy. AI literacy includes understanding how AI systems function, evaluating their limitations and bias, using AI ethically, and applying these tools in academic and real-world contexts. Research shows that students across grade levels can develop accurate understandings of AI when instruction is thoughtfully designed (Druga et al., 2019; DiPaola & Gillespie, 2021).

AI LITERACY

Maryland's approach aligns with national and international guidance that frames AI literacy as a core component of digital citizenship, critical thinking, and future readiness (UNESCO, 2023; U.S. Department of Education, 2023).

AI literacy instruction in Maryland emphasizes understanding how data and design choices shape AI behavior, evaluating AI-generated content for accuracy and bias, recognizing appropriate versus inappropriate uses, and considering ethical, and societal impacts. These expectations align with established frameworks, including the AI4K12 Five Big Ideas, the CSTA K–12 Computer Science Standards, and statewide digital citizenship, media literacy, and content standards. AI literacy is not confined to computer science but integrated across disciplines, including ELA, mathematics, science, social studies, CTE, and the arts.

AI literacy equips students to evaluate information systems, recognize potential harms, and make informed judgments about appropriate use, competencies increasingly essential for postsecondary success and workforce participation (Long & Magerko, 2020; World Economic Forum, 2023).

Maryland's AI literacy competencies progress across grade bands, from early awareness in the primary grades to critical evaluation, ethical reasoning, and domain-specific application in high school. Students in grades 2–12 are expected to engage annually in AI literacy instruction, recognizing that repeated exposure builds durable understanding (Hobbs, 2021). Instruction may occur through core coursework, digital citizenship lessons, computer science and CTE programs, or interdisciplinary experiences, allowing local flexibility within statewide expectations.

CURRICULUM INTEGRATION

Educator capacity is essential to effective AI literacy. Maryland will support instruction through model lessons, curated resources, professional learning for educators across disciplines, and guidance on assessing AI literacy while maintaining academic integrity (Zhai et al., 2023).

Ethical use is foundational. Students must learn to protect privacy, recognize harmful or misleading outputs, understand the limits of machine reasoning, and use AI responsibly (Floridi & Cows, 2019). The goal is to prepare every Maryland student to engage with AI systems confidently, critically, and ethically, supporting success in college, careers, and life in an AI-mediated world.



Operational and Administrative Uses

Artificial intelligence is increasingly used in school system operations to support efficiency, communication, and decision making. Because these systems operate at scale and can affect student access to services and opportunities, they present heightened risks related to privacy, equity, transparency, and accountability.

Maryland treats operational AI as high stakes use case requiring caution, transparency, and continuous human oversight. Research shows that automated systems can reinforce inequities, obscure decision logic, and encourage over-reliance if not carefully governed (Eubanks, 2018; Obermeyer et al., 2019)

All AI-assisted operational decisions that affect student opportunities must undergo human review. Educators and administrators retain authority to interpret, contextualize, and override AI recommendations and must be trained to recognize model limitations and bias.

Districts must communicate clearly with families about operational AI use, evaluate systems for bias before adoption, monitor performance and impacts over time, limit use to approved educational purposes, and maintain procedures for addressing inaccurate or harmful outputs.

Maryland will support districts through model policies, guidance, professional learning, vetted resources, and reporting mechanisms to ensure operational AI strengthens equitable, transparent, and student-centered school system functioning.

Operational AI includes tools that rely on predictive modeling, classification, or automation to support functions such as attendance and early-warning systems, enrollment forecasting, scheduling, transportation, automated communication, and security monitoring.




Procurement and Vendor Governance

Artificial intelligence tools used in Maryland schools must meet high standards for safety, transparency, privacy, equity, and educational value. Because AI products vary widely in data practices, accuracy, and risk, procurement processes must ensure that tools support learning and operations without compromising student rights or causing harm (Kumar et al., 2022).

Maryland requires rigorous vendor vetting, strong contractual protections, documented risk assessments, and ongoing monitoring. Procurement decisions must align with state goals, legal requirements, equity priorities, and responsible data stewardship, consistent with federal guidance on ethical AI adoption in education (U.S. Department of Education, 2023).

Districts must prioritize tools with clear educational or operational value, transparent data practices, compliance with FERPA, COPPA, state privacy laws, and civil rights protections, and demonstrated attention to bias, fairness, and accessibility. Approval is an ongoing process, not a one-time decision. To support consistent practice, Maryland will provide vendor vetting rubrics, model contract language, a centralized list of vetted tools, professional learning, technical assistance, and reporting mechanisms to ensure statewide AI procurement protects students and upholds equity, transparency, and legal compliance.



Conclusion

Artificial intelligence will continue to shape education and society. Maryland's goal is not to chase the latest technologies, but to ensure that AI deepens learning, advances equity, protects student wellbeing, and supports educator judgment. AI does not replace human relationships or professional expertise; it serves as a tool in support of Maryland's educational mission.

This guidance establishes a statewide foundation for responsible AI use grounded in transparency, privacy, equity, and continuous improvement. It affirms the role of educators, clarifies shared responsibilities, and emphasizes thoughtful adoption informed by evidence and risk.

Responsible AI use is ongoing. As technologies evolve, Maryland will adapt policies, support local implementation, and engage educators, students, and families.

Maryland's guiding test is simple: does AI help students learn and thrive? If yes, it is supported. If not, Maryland will proceed with caution—centering humanity, safety, and intellectual rigor in every decision

References

- Benjamin, R. (2019). Race after technology: Abolitionist tools for the New Jim Code. Polity.
- CAST. (2018). Universal Design for Learning guidelines version 2.2. CAST. <https://udlguidelines.cast.org>
- Crawford, K. (2021). The atlas of AI: Power, politics, and the planetary costs of artificial intelligence. Yale University Press.
- DiPaola, D., & Gillespie, R. (2021). Young children's understanding of artificial intelligence. *Computers & Education*, 168, 104212. <https://doi.org/10.1016/j.compedu.2021.104212>
- Druga, S., Williams, R., Breazeal, C., & Resnick, M. (2019). How children understand AI: Removal, agency, and the illusion of control. *Proceedings of the ACM Interaction Design and Children Conference*, 123–133. <https://doi.org/10.1145/3311927.3323157>
- Eubanks, V. (2018). Automating inequality: How high-tech tools profile, police, and punish the poor. St. Martin's Press.
- Floridi, L., & Cowls, J. (2019). A unified framework of five principles for AI in society. *Harvard Data Science Review*, 1(1). <https://doi.org/10.1162/99608f92.8cd550d1>
- Goddard, K., Roudsari, A., & Wyatt, J. C. (2012). Automation bias: A systematic review. *Journal of the American Medical Informatics Association*, 19(1), 121–127. <https://doi.org/10.1136/amiajnl-2011-000089>
- Hobbs, R. (2021). Media literacy in a digital world (3rd ed.). Wiley.
- Holmes, W., Bialik, M., & Fadel, C. (2019). Artificial intelligence in education: Promises and implications for teaching and learning. Center for Curriculum Redesign.
- Holmes, W., Porayska-Pomsta, K., Holstein, K., Sutherland, E., Baker, T., Shum, S. B., Rodrigo, M. T., Cukurova, M., Bittencourt, I. I., & Koedinger, K. R. (2021). Ethics of AI in education: Towards a community-wide framework. *International Journal of Artificial Intelligence in Education*. <https://doi.org/10.1007/s40593-021-00239-1>
- Koehn, P., & Knowles, R. (2017). Six challenges for neural machine translation. *Proceedings of the First Workshop on Neural Machine Translation*, 28–39. <https://doi.org/10.18653/v1/W17-3204>
- Kumar, P., Vitak, J., & Zimmer, M. (2022). Privacy and security concerns with educational technology: A student perspective. *Computers & Education*, 176, 104353. <https://doi.org/10.1016/j.compedu.2021.104353>
- Liang, P. P., Wu, Z., Zosso, N., Wu, J., & He, X. (2023). GPT detectors are biased against non-native English writers. *arXiv*. <https://doi.org/10.48550/arXiv.2304.02819>
- Long, D., & Magerko, B. (2020). What is AI literacy? *Proceedings of the CHI Conference on Human Factors in Computing Systems*, 1–16. <https://doi.org/10.1145/3313831.3376727>
- Noble, S. U. (2018). Algorithms of oppression: How search engines reinforce racism. NYU Press. <https://doi.org/10.18574/nyu/9781479833641.001.0001>
- Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage population health. *Science*, 366(6464), 447–453. <https://doi.org/10.1126/science.aax2342>
- Reddy, S., Allan, M., Coghlan, S., & Cooper, A. (2020). A governance model for the application of AI in health care. *Journal of the American Medical Informatics Association*, 27(3), 491–497. <https://doi.org/10.1093/jamia/ocz217>
- Schiff, D. (2022). Education for AI, not AI for education: The role of ethics in national AI policy. *International Journal of Artificial Intelligence in Education*, 32(3), 527–563. <https://doi.org/10.1007/s40593-021-00270-2>
- UNESCO. (2023). Guidance for generative AI in education and research. UNESCO.
- U.S. Department of Education, Office for Civil Rights, & Office of Educational Technology. (2024). Avoiding the discriminatory use of artificial intelligence in education. U.S. Department of Education.
- U.S. Department of Education, Office of Educational Technology. (2023). Artificial intelligence and the future of teaching and learning: Insights and recommendations.
- World Economic Forum. (2023). The future of jobs report. WEF.
- Zhai, X., Zhang, Z., & Cui, Y. (2023). Factors influencing teacher adoption of AI in classrooms. *Educational Technology Research and Development*, 71, 421–445. <https://doi.org/10.1007/s11423-022-10>

